

A parallel iterated local search for the traveling salesman problem as a multi-core solution using OpenMP

Isaac Kosloski Oliveira¹

¹Faculdade de Computação

Universidade Federal de Mato Grosso do Sul (UFMS) Campo Grande – MS – Brazil

isaac.kosloski@ufms.br

Abstract. *This write-up reports the results of an implementation of the Iterated Local Search algorithm (ILS), both sequential and parallel, for the Traveling Salesman Problem (TSP). The results are disposable in two phases. The first phase compares the evaluation measures of ILS on 1 and 12 cores (sequential and parallel, respectively, and optimal). The second phase compares the proposed parallel algorithm with the reported results of metaheuristics algorithms that were used to solve the TSP in the literature.*

1. Introduction

One of the most prominent components in the set of combinatorial optimization problems, dating back to the 19th century, is certainly the *traveling salesman problem* (TSP). This job of finding the minimum weight cycle in a given graph is one of the few mathematical problems that frequently come about in the most popular scientific press [Reinelt 1994].

2. Literature Review

2.1. Solving Traveling Salesman Problem Using Parallel River Formation Dynamics Optimization Algorithm on Multi-core Architecture Using Apache Spark

Year: 2024

It uses the parallel RFD algorithm for solving the TSP, comparing speedup, running time, and efficiency on 1 (sequential), 4, 8, and 16 cores. Then, compare to three parallel water-based algorithms (*Water Flow*, *Intelligent Water Drops* and *Water Cycle*). Then, the proposed algorithm will be compared with the reported metaheuristics used to solve TSP in the literature.

Similar benchmarks:

- d198;
- lin318;
- rat783.

Can be used for comparison:

- Distance;
- Accuracy;
- Running time;
- Speedup.

2.2. The AddACO: A bio-inspired modified version of the ant colony optimization algorithm to solve travel salesman problems

Year: 2024

It's a proposal to solve the TSP with the AddACO algorithm (it's a version of the Ant Colony Optimization method characterized by a modified probabilistic law at the basis of the exploratory movement of the artificial insects). In particular, the ant decisional rule is here set to the amount in a linear convex combination of competing behavioral stimuli. It has an additive form (hence the name of our algorithm) rather than the canonical multiplicative one. The AddACO intends to address two conceptual shortcomings that characterize classical ACO methods:

- (i) the population of artificial insects is, in principle, allowed to simultaneously minimize/maximize all migratory guidance cues (which is implausible from a biological/ecological point of view).
- (i) a given edge of the graph has a null probability of being explored if at least one of the movement traits is equal to zero, i.e., regardless of the intensity of the others (this, in principle, reduces the exploratory potential of the ant, colony).

Similar benchmarks:

- lin318;

Can be used for comparison:

→ Distance;

2.3. A novel hybrid swarm intelligence algorithm for solving TSP and desired-path-based online obstacle avoidance strategy for AUV

Year: 2024

Similar benchmarks:

- none;

Can be used for comparison:

→ Distance;

2.4. Discrete artificial bee colony algorithm with fixed neighborhood search for traveling salesman problem

Year: 2024

It proposes a discrete artificial bee colony algorithm with a fixed neighborhood search for the traveling salesman problem (TSP) called DABC-FNS. The solution obtained by the algorithm is expressed by a positive integer coding method. Meanwhile, the local enhancement strategy and the 2-opt strategy with fixed neighborhood search are introduced to improve the ABC algorithm's solution accuracy.

Similar benchmarks:

- d198;
- a280;
- rat783.

Can be used for comparison:

→ Distance;

2.5. The Discrete Carnivorous Plant Algorithm with Similarity Elimination Applied to the Traveling Salesman Problem

Year: 2022

It uses a combination of six steps: first, the algorithm redefines subtraction, multiplication, and addition operations, which aims to ensure that it can switch from continuous space to discrete space without losing information; second, a simple sorting grouping method is proposed to reduce the chance of being trapped in a local optimum; third, the similarity-eliminating operation is added, which helps to maintain population diversity; fourth, an adaptive attraction probability is proposed to balance exploration and the exploitation ability; fifth, an iterative local search (ILS) strategy is employed, which is beneficial to increase the searching precision; finally, to evaluate its performance, DCPA is compared with nine algorithms.

Similar benchmarks:

- d198;
- lin318;
- pcb442;
- rat783;
- fl1577.

Can be used for comparison:

→ Distance;

3. Problem description

The traveling salesman problem is one of the most well-known problems in combinatorial optimization and a member of the NP-hard problems [Esra'a Alhenawi and Hussien 2024]. Close related to the Hamiltonian cycle problem, its applications can also be modeled as a graph problem.

Modeling this problem as a complete graph with n vertices, where the set of the vertices represents a group of cities, and the set of the arcs typifies a group of roads interconnecting the cities. The main target is for the salesman to make a tour (hamiltonian cycle), visiting each city exactly once and finishing at the city he has started from [Thomas H. Cormen and Stein 2009].

In the standard classical problem, the salesman incurs a nonnegative cost $c(i, j)$ to travel from city i to city j . The desired total cost of this tour should be the minimum distance, where the total cost is the sum of the individual costs along the edges of the tour, i.e., the weight minimum cycle in the graph problem.

In a more formal specification for the TSP, given a complete graph $K_n = (V, E)$, where V is the finite set of vertices (typifying the cities), and E the set of edges (typifying the roads), that has nonnegative cost $c(v_i, v_j)$ (typifying the distance between two cities) associated with each edge $(v_i, v_j) \in E$. An acceptable solution is a permutation σ of V , represented as $S_\sigma = (v_{\sigma(1)}, v_{\sigma(2)}, \dots, v_{\sigma(n)})$, where $n = |V|$, which minimizes the cycle cost $C(\sigma)$, then:

$$\min C(\sigma) = \sum_{i=0}^{n-1} c(v_i, v_{i+1}) + c(v_{n-1}, v_0) \quad (1)$$

Given S_n , the set of all symmetric permutations in V of n elements, the optimal solution S_{opt} is a minimum weight Hamiltonian cycle.




Figure 1. K_5 graph for example 1.

For example, given the graph K_5 in figure 1, an solution is the cycle $(v_1, v_2, v_5, v_3, v_4, v_1)$, illustrated in figure 2, which $C = c(v_1, v_2) + c(v_2, v_5) + c(v_5, v_3) + c(v_3, v_4) + c(v_4, v_1)$, as described in 1.




Figure 2. An instance of the traveling-salesman problem, on K_5 graph, for example 1.

Now, we can define the TSP as an optimization problem more carefully for implementation. Let a variable $x_{ij} \in \mathbb{B}$ for each edge $(v_i, v_j) \in E$, i.e. for $i = 1, 2, \dots, n$ and $j = 1, 2, \dots, n$, where $i \neq j$. Let a variable $c_{ij} \in \mathbb{R}$ be a cost (weight) value associated with each edge $(v_i, v_j) \in E$. With those statements, the objective function is formulated, and we can elucidate the constraints for completing the formulation. We desire a Hamiltonian cycle, i.e., include each vertex in our set exactly one time, then we must let one edge as a way to "get in" and another to "get out", as the variable x_{ij} already is related to an edge, and for each vertex, there are exactly 2 edges associated, then we just set the edges related to the vertex as 1. Other constraints are related to sub-cycles, so an alternative could be to let a set S of all sub-sets from V in which the number of elements is between 2 and $n - 2$. For each $S_\sigma \in S$, we can set at least one edge of the selected cycle to "escape" of this set S . Given these statements, we can formulate:

$$z = \min \sum_{i=1}^n \sum_{j=1}^n c_{ij} X_{ij} :$$

$$\sum_{i=1}^n X_{ij} = 1,$$

$$\sum_{j=1}^n X_{ij} = 1,$$

$$\sum_{i \in S} x_{ij} \geq 1, \forall S_\sigma \in S,$$

$$x_{ij} \in \mathbb{B}, i \neq j.$$

4. Iterated Local Search

The *Iterated Local Search* (ILS) is a sophisticated metaheuristics algorithm designed to solve optimization problems. It operates by iteratively generating embedded solutions and comparing them to provide an acceptable one.

Imagine an optimization heuristic algorithm, a local search, tailored for a specific problem. This can be implemented as a *LocalSearch* procedure. The question that arises is, 'Can we iteratively optimize?' If the answer is 'yes', then the optimization obtained is not just significant, but invaluable. To comprehend the workings of the local search, we assume it to be deterministic and memoryless.

Let C be the cost function of our optimization problem, a function that we strive to minimize. Let S be the finite set of all solutions s . In figure 4, a high-level block diagram of local search is illustrated, where given an input s , we always generate the same output s^* with a cost less than or equal to $C(s)$. The *LocalSearch* defines a n to 1 mapping from set S to a smaller set S^* with optimal local solutions s^* . The main feature of a local search is a neighbor structure; from this, we can understand S as some topological structure, not just a set, where this allows agile to go from one solution to another [Helena R. Lorenço and STUTZLE].



Figure 3. Block diagram of LocalSearch procedure




Figure 4. Block diagram of Iterated Local Search procedure

Algorithm 1 Iterated Local Search

```
PROCEDURE Iterated Local Search
   $s_0$  = GenerateInitialSolution()
   $s^*$  = LocalSearch( $s_0$ )
  REPEAT
     $s'$  = Perturbation( $s^*$ , history)
     $s^{*'} = \text{LocalSearch}(s')$ 
     $s^*$  = AcceptanceCriterion( $s^*$ ,  $s^{*'}$ , history)
  UNTIL termination condition met
END
```

5. Implementing the TSP Using ILS

6. Experiment Setup

6.1. Implement Language and Frameworks

The implementation uses C++20 and OpenMP 5.2, compiled using GCC 12.2.0.

6.2. Testing Enviroment

The code was compiled and run on a machine x86_64, AMD Ryzen 5 5600GT with Radeon Graphics, 6 cores, and 12 threads. With Debian 12.2.0-14. The program was compiled using `g++ component.cpp node.cpp scanner.cpp functions.cpp mainExec.cpp -o ../bin/TspPar -fopenmp -Wall -pedantic`, and tested using a shell script.

6.3. Datasets - Inputs

For the datasets, nine TSP benchmarks were downloaded from TSPLIB [TspLib], which includes d198, a280, lin318, pcb442, rat783, u1060, pcb1173, d1291, and fl1577 have been used in this write-up for evaluating the proposed algorithm performance. The selected benchmarks varied in several cities where each city was represented by a 2D-Euclidian coordinate. The benchmarks' input was a *.tsp* file, with the format:

```
NAME : <benchmark name>
COMMENT : <benchmark description>
TYPE : TSP
DIMENSION : <dimension>
EDGE_WEIGHT_TYPE : EUC_2D
NODE_COORD_SECTION
<i> <coordinate  $x_i$ > <coordinate  $y_i$ >
EOF
```

The table 1 displays benchmarks properties.

| Benchmark name | Dimension | Description |
|----------------|-----------|---|
| d198.tsp | 198 city | Drilling problem (Reinelt) |
| a280.tsp | 280 city | drilling problem (Ludwig) |
| lin318.tsp | 318 city | 318-city problem (Lin/Kernighan) |
| pcb442.tsp | 442 city | Drilling problem (Groetschel/Juenger/Reinelt) |
| rat783.tsp | 783 city | Rattled grid (Pulleyblank) |
| u1060.tsp | 1060 city | Drilling problem problem (Reinelt) |
| pcb1173.tsp | 1173 city | Drilling problem (Juenger/Reinelt) |
| d1291.tsp | 1291 city | Drilling problem (Reinelt) |
| fl1577.tsp | 1577 city | Drilling problem (Reinelt) |

Table 1. TSP benchmarks' properties

6.4. Data Analysis - outputs

The output data, *.sol* file, has the format:

Iterations: 1000

Time: <time in seconds> sec - <time in minutes> min - <time in hours> horas
 <fraction of time in hours> h <fraction of time in minutes> min

Problem dimension: <dimension>

Total distance: <total path distance>

[v_1] [v_2] [v_3] ... [v_i]

Where v_i it's the city in the position i on the tour sequence. For each benchmark, there are 30 different output *.sol* files, to define the trial result statistically accurately.

6.5. Evaluation Measures

This section presents the evaluation metrics that are used for evaluating the proposed method.

1. Distance: the value of the best route found.
2. Accuracy: the percentage of retrieving the best route correctly. The optimal solution is provided in TSPLIB [TspLib].
3. Running time: is the duration between the end and the beginning of the main part (ILS algorithm itself) of the program running, using **chrono::high_resolution_clock::now()** for the sequential program and **omp_get_wtime()** for the parallel program.

$$RT = \text{finish_time} - \text{start_time} \quad (2)$$

4. Speedup: is the improvement in speed of execution of a task executed on two similar architectures with different resources:

$$SP = \frac{T_s}{T_p(n)} \quad (3)$$

- T_s : Sequential running time;

- $T_p(n)$: Parallel running time in function of n ;
- n : Number of cores.

5. Efficiency: represents the speedup divided by the number of cores

$$EF = \frac{SP}{n} \quad (4)$$

7. Experimental results

The results are taken by the average of 30 trials on each instance.

7.1. Comparison of Sequential and Parallel ILS

Table 2. Average distance of each symmetric TSP benchmark, the Sequential and Parallel ILS, then the optimal solution.

| Instance | Sequential Distance | Parallel Distance | Optimal |
|----------|---------------------|---------------------|----------------------|
| d198 | 1.628×10^4 | 1.721×10^4 | 1.578×10^4 |
| a280 | 2.643×10^3 | 2.729×10^3 | 2.579×10^3 |
| lin318 | 4.382×10^4 | 4.763×10^4 | 4.203×10^4 |
| pcb442 | 5.172×10^4 | 5.348×10^4 | 5.078×10^4 |
| rat783 | 9.271×10^3 | 9.659×10^3 | 8.806×10^3 |
| u1060 | 2.467×10^5 | 2.416×10^5 | 2.241×10^5 |
| pcb1173 | 0.000 | 6.498×10^4 | 5.689×10^4 |
| d1291 | 0.000 | 5.690×10^4 | 5.080×10^4 |
| fl1577 | 0.000 | 2.966×10^4 | 2.2249×10^4 |

Table 3. Accuracy of each symmetric TSP benchmark for Sequential and Parallel ILS.

| Instance | Sequential Accuracy (%) | Parallel Accuracy (%) |
|----------|-------------------------|-----------------------|
| d198 | 96.91 | 91.97 |
| a280 | 97.58 | 94.50 |
| lin318 | 95.91 | 88.25 |
| pcb442 | 98.18 | 94.95 |
| rat783 | 94.98 | 91.17 |
| u1060 | 90.84 | 92.75 |
| pcb1173 | 00.00 | 87.55 |
| d1291 | 00.00 | 89.27 |
| fl1577 | 00.00 | 75.02 |

Table 4. Running time, speedup, and efficiency of each symmetric TSP benchmark for Sequential and Parallel ILS.

| Instance | Sequential Time (s) | Parallel Time (s) | Speedup | Efficiency (%) |
|----------|---------------------|---------------------|---------|----------------|
| d198 | 6.114×10^1 | 8.649 | 7.069 | 58.91 |
| a280 | 1.755×10^2 | 2.411×10^1 | 7.279 | 60.66 |
| lin318 | 2.775×10^2 | 3.941×10^1 | 7.041 | 58.68 |
| pcb442 | 7.463×10^2 | 1.018×10^2 | 7.331 | 61.09 |
| rat783 | 4.620×10^3 | 5.999×10^2 | 7.701 | 64.18 |
| u1060 | 6.542×10^2 | 5.062×10^3 | 2.404 | 20.03 |
| pcb1173 | 0.000 | 2.183×10^3 | 0.000 | 00.00 |
| d1291 | 0.000 | 2.712×10^3 | 0.000 | 00.00 |
| fl1577 | 0.000 | 5.797×10^3 | 0.000 | 00.00 |

References

- Esra'a Alhenawi, Ruba Abu Khurma, R. D. and Hussien, A. G. (2024). Solving traveling salesman problem using parallel river formation dynamics optimization algorithm on multi-core architecture using apache spark.
- Helena R. Lorenzo, O. M. and STUTZLE, T. Iterated local search.
- Reinelt, G. (1994). *The Traveling Salesman: Computational Solutions for TSP Applications*. Springer-Verlag Berlin Heidelberg, 1^o edition.
- Thomas H. Cormen, Charles E. Leiserson, R. L. R. and Stein, C. (2009). *Introduction to Algorithms*. The MIT Press, 3^o edition.
- TspLib. Disponível em: <http://comopt.ifl.uni-heidelberg.de/software/TSPLIB95/>. Acesso em: 11 de setembro 2023.