

A parallel iterated local search and Greedy Randomized Adaptive Search Procedure for the traveling salesman problem as a multi-core solution using OpenMP

Isaac Kosloski Oliveira¹

¹Faculdade de Computação

Universidade Federal de Mato Grosso do Sul (UFMS) Campo Grande – MS – Brazil

isaac.kosloski@ufms.br

Abstract. *This write-up reports the results of an implementation of the Iterated Local Search algorithm (ILS), both sequential and parallel, for the Traveling Salesman Problem (TSP). The results are disposable in two phases. The first phase compares the evaluation measures of ILS on 1 and 12 cores (sequential and parallel, respectively, and optimal). The second phase compares the proposed parallel algorithm with the reported results of metaheuristics algorithms that were used to solve the TSP in the literature.*


1. Introduction

One of the most prominent components in the set of combinatorial optimization problems, dating back to the 19th century, is certainly the *traveling salesman problem* (TSP). Finding the minimum weight cycle in a given graph is one of the few mathematical problems frequently occurring in the most popular scientific press [Reinelt 1994].

Certainly, practical applications of the TSP can be found in the real world. A standard symmetric traveling salesman problem application is the drilling problem for printed circuit boards (PCBs). To connect conductors in different layers or position the pins of integrated circuits (ICs), holes must be drilled through the board. In most PCBs, those holes may be of different diameters. To avoid having to drill two holes of different diameters consecutively, which can consume an amount of undesirable time, it is clear at the outset that one has to choose some diameter, drill all holes, and do the same consecutively. This problem can be interpreted as a sequence of symmetric traveling salesman problems [Reinelt 1994].

Another problem that can be modeled as a symmetric TSP is the overhauling of gas turbine engines, which occurs when aircraft's gas turbine engines need to be overhauled. There are nozzle-guide vane assemblies at each turbine stage to guarantee a uniform gas flow through the turbines. Such an assembly consists of a number of nozzle guide vanes affixed about its circumference. All these vanes have individual characteristics, and the correct placement of the vanes can result in substantial benefits for the plant, making the system more stable, like reducing vibration, increasing uniformity of flow, and reducing fuel consumption, which helps control engineering too.

The task of determining a sequence of operations for the control of a robot that leads to the shortest overall processing time is similar to finding the shortest Hamiltonian



path, given that this application's difficulty arises because precedence constraints must be observed. This problem can be formulated as a TSP variation by applying similar methods to solve the problem. Different problems can be implemented as a TSP, like X-ray crystallography problems, the order-picking problem in warehouses, computer wiring, clustering of a data array, seriation in archeology, vehicle routing problems, scheduling problems, and mask plotting in PCS production, among others.

The TSP incurs the cost of finding a minimum weight Hamilton cycle. In graph theory, it means finding a cycle that contains every vertex of a given graph. A graph is *hamiltonian* if it contains a Hamilton cycle. For this problem, there is no nontrivial necessary and sufficient condition for a graph to be Hamiltonian, and the problem of finding such a condition is one of the main unsolved problems of graph theory [Bondy and Murty 1976].


The expected graph, where one should search for a Hamilton cycle, is a weighted complete graph. Such a cycle should be comprehended as an optimal cycle. In contrast with other classical problems in graph theory, like the shortest path problem or the connector problem, no exact algorithm for solving the traveling salesman problem is known.

Since the TSP is largely applicable, it is desirable to have a method for obtaining a reasonably good solution, which is not necessarily optimal [Bondy and Murty 1976]. Hence, approximation algorithms and metaheuristics are used to solve the TSP.

With the advent of the theory of NP-completeness, approximation algorithms and metaheuristics became more prominent as the need to generate near-optimal solutions for NP-hard optimization problems and deal with computational intractability. It is possible for some problems to generate near-optimal solutions quickly, whereas for others, generating probably good suboptimal solutions is as difficult as generating optimal ones. Approximation algorithms and metaheuristics have been developed to solve various problems, including the TSP. These algorithms have theoretical value since their time complexity is a high-order polynomial or a huge constant associated with the time complexity bound, which provides real-world computational time feasibility [].

A simple approach is preferable in designing metaheuristics algorithms, both in concept and application. Iterated Local Search (ILS) can be introduced with this in mind. The main procedure can be interpreted as an embedded heuristic called *LocalSearch*. Wichi provides a local optimal solution, given a neighborhood structure. To generate better solutions, one can iterate the *LocalSearch* as much as necessary and slightly modify the solution so that the heuristic search for new solutions in a different neighborhood search space [Helena R. Lorenço and STUTZLE].

Another approach, as simple as the ILS, is the Greedy Randomized Adaptive Search Procedure (GRASP). Where a *LocalSearch* exists too, they also differ in constructing the solution that the *LocalSearch* should look for. The *LocalSearch* algorithm works iteratively, replacing the current solution with a better solution in the neighborhood of the current solution. The solution that the *LocalSearch* receives is generated by a *GreedyRandomizedConstructor*, where given a random point in the previous solution, one selects the closest neighbors to build a restricted candidate list (RCL). Our solution is constructed by selecting one of those elements at random. Hence, successively applying this RCL build and solution construction gives a relevant solution for the *LocalSearch* returns an



optimized solution. Again, one could iterate these steps to generate better solutions.

Until now, the main issue was the quality of the solutions in a feasible running time. But going beyond this, concurrency and parallelism could be applied to get a better time. Improving the running time and guaranteeing the same quality of solutions could scale the difficulty of implementation. The first approach to consider is a multi-core implementation, letting the metaheuristics explore parallelism in CPU cores. The framework OpenMP covers user-directed parallelization, which provides high-architecture code-writing, but the user is still free to implement concurrency and parallelism as needed [].

2. Literature Review

Research on optimization metaheuristics to design solutions for solving TSP problems has been developed since the early century, but it is still a subject of study.

A recent solution [Esra'a Alhenawi and Hussien 2024] uses the parallel RFD (*River Formation Dynamics*) algorithm for solving the TSP, comparing speedup, running time, and efficiency on sequential code and parallel, with 4, 8, and 16 cores. Then, compare to three parallel water-based algorithms (*Water Flow*, *Intelligent Water Drops* and *Water Cycle*). Then, the proposed algorithm will be compared with the reported metaheuristics used to solve TSP in the literature. It brings an extensive evaluation measure set like distance, accuracy, running time, and speedup.

Another recent proposal [Scianna 2024] is to solve the TSP with the AddACO algorithm (a version of the Ant Colony Optimization method characterized by a modified probabilistic law at the basis of the exploratory movement of the artificial insects). In particular, the ant decisional rule is set to the amount in a linear convex combination of competing behavioral stimuli. It has an additive form (hence the algorithm's name) rather than the canonical multiplicative one. The AddACO intends to address two conceptual shortcomings that characterize classical ACO methods:

- (i) the population of artificial insects is, in principle, allowed to simultaneously minimize/maximize all migratory guidance cues (which is implausible from a biological/ecological point of view).
- (i) a given edge of the graph has a null probability of being explored if at least one of the movement traits is equal to zero, i.e., regardless of the intensity of the others (this, in principle, reduces the exploratory potential of the ant colony).

A proposal of a discrete artificial bee colony algorithm with a fixed neighborhood search for the traveling salesman problem (TSP) called DABC-FNS [Xing Li and Shao 2024], where the solution obtained by the algorithm is expressed by a positive integer coding method. Meanwhile, the local enhancement strategy and the 2-opt strategy with fixed neighborhood search are introduced to improve the ABC algorithm's solution accuracy.

The DCPA (*Discrete Carnivorous Plant Algorithm*) with Similarity Elimination Applied to the Traveling Salesman Problem is another idea [Pan-Li Zhang and Zhang 2022], where it uses a combination of six steps: first, the algorithm redefines subtraction, multiplication, and addition operations, which aims to ensure that it can switch from continuous space to discrete space without losing information; second, a simple sorting grouping method is proposed to reduce the chance of being trapped in a local optimum; third, the similarity-eliminating operation is added, which helps to maintain population

diversity; fourth, an adaptive attraction probability is proposed to balance exploration and the exploitation ability; fifth, an iterative local search (ILS) strategy is employed, which is beneficial to increase the searching precision; finally, to evaluate its performance, DCPA is compared with nine algorithms.

3. Problem description

The traveling salesman problem is one of the most well-known problems in combinatorial optimization and a member of the NP-hard class of problems [Esra'a Alhenawi and Hussien 2024]. Close related to the Hamiltonian cycle problem, its applications can also be modeled as a graph problem.

Modeling this problem as a complete graph with n vertices, where the set of the vertices represents a group of cities, and the set of the arcs typifies a group of roads interconnecting the cities, the main target is for the salesman to make a tour (hamiltonian cycle), visiting each city exactly once and finishing at the city he has started from [Thomas H. Cormen and Stein 2009].

In the standard classical problem, the salesman incurs a nonnegative cost $c(i, j)$ to travel from city i to city j . The desired total cost of this tour should be the minimum distance, where the total cost is the sum of the individual costs along the edges of the tour, i.e., a minimum weight cycle in the graph problem.

In a more formal specification for the TSP, given a complete graph $K_n = (V, E)$, where V is the finite set of vertices (representing the cities), and E the set of edges (representing the roads), that has nonnegative cost $c(v_i, v_j)$ (representing the distance between two cities) associated with each edge $(v_i, v_j) \in E$. An acceptable solution is a permutation σ of V , represented as $S_\sigma = (v_{\sigma(1)}, v_{\sigma(2)}, \dots, v_{\sigma(n)})$, where $n = |V|$, which minimizes the cycle cost $C(\sigma)$, then:

$$\min C(\sigma) = \sum_{i=0}^{n-1} c(v_i, v_{i+1}) + c(v_{n-1}, v_0) \quad (1)$$

Given S_n , the set of all symmetric permutations in V of n elements, the optimal solution S_{opt} is a minimum weight Hamiltonian cycle.





Figure 1. K_5 graph for example 1.



For example, given the graph K_5 in Figure 1, one possible solution is the cycle $(v_1, v_2, v_5, v_3, v_4, v_1)$, illustrated in Figure 2, with cost $C = c(v_1, v_2) + c(v_2, v_5) + c(v_5, v_3) + c(v_3, v_4) + c(v_4, v_1)$, as described in Equation (1).




Figure 2. An instance of the traveling-salesman problem, on K_5 graph, for example 1.

Now, we can define the TSP as an optimization problem more carefully for implementation. Let a variable $x_{ij} \in \{0, 1\}$ for each edge $(v_i, v_j) \in E$, i.e. for $i = 1, 2, \dots, n$ and $j = 1, 2, \dots, n$, where $i \neq j$. Let a variable $c_{ij} \in \mathbb{R}$ be a cost (weight) value associated with each edge $(v_i, v_j) \in E$. With those statements, the objective function is formulated in (2a), where it minimizes the sum of all edges x_{ij} weighted by a cost c_{ij} associated with each edge. The constraints can be elucidated to complete the formulation. We desire a Hamiltonian cycle, i.e., include each vertex in our set exactly once. We must let one edge as a way to “get in” and another to “get out”, as the variable x_{ij} already is related to an edge, and for each vertex, there are exactly two edges associated, as in (2c) and (2d). We just set the edges related to the vertex as 1. Other constraints are related to subcycles, so an alternative could be to allow a set S of all subsets from V in which the number of elements is between 2 and $n - 2$. For each $S_\sigma \in S$, we can set at least one edge of the selected cycle to “escape” of this set S , as in (2d).

$$z = \min \sum_{i=1}^n \sum_{j=1}^n c_{ij} x_{ij} : \quad (2a)$$

$$\sum_{i=1}^n x_{ij} = 1, \quad (2b)$$

$$\sum_{j=1}^n x_{ij} = 1, \quad (2c)$$

$$\sum_{i \in S} x_{ij} \geq 1, \forall S_\sigma \in S, \quad (2d)$$

$$x_{ij} \in \{0, 1\}, i \neq j.$$

4. Iterated Local Search

The *Iterated Local Search* (ILS) is a sophisticated metaheuristics algorithm designed to solve optimization problems. It operates by iteratively generating embedded solutions and comparing them to provide an acceptable one.

Imagine an optimization heuristic algorithm, a local search, tailored for a specific problem. This can be implemented as a *LocalSearch* procedure. The question that arises is, ‘Can we iteratively optimize?’ If the answer is ‘yes’, then the optimization obtained is significant and invaluable. To comprehend the workings of the local search, we assume it to be deterministic and memoryless.

Let C be the cost function of our optimization problem, a function that we strive to minimize. Let S be the finite set of all solutions s . In Figure 3, a high-level block diagram of local search is illustrated where, given an input s , we always generate the same output s^* with a cost less than or equal to $C(s)$. The *LocalSearch* defines a n to 1 mapping from set S to a smaller set S^* with optimal local solutions s^* . The main feature of a local search is a neighbor structure; from this, we can understand S as some topological structure, not just a set, where this allows going from one solution to another [Helena R. Lorenzo and STUTZLE].

Hence, a helpful concept is the *basin of attraction* of a local minimum s^* , a set of s mapped to s^* under the local search routine. The *LocalSearch* is provided with a solution s that gives a local minimum s^* in the *basin of attraction*. To explore solutions in S out of the local *basin of attraction*, we can apply a change, or *Perturbation*, that leads to an intermediate state s' , then, the *LocalSearch* can be applied. A new minimum local solution $s^{*'}$ is given, as represented in the second part of the block diagram in Figure 4.

The next step is, given the two local minimum solutions, i.e., s^* and $s^{*'}$, we must establish a criterion that decides which one our solution should be. The *AcceptanceCriterion* implements this task, choosing which solution is better for the next step. It can be defined according to the problem and its desired solution, e.g., the less-cost solution or the more diversified solution. Once the base structure is defined, the block diagram in Figure 4 provides a view of how the ILS metaheuristics work.



Figure 3. Block diagram of LocalSearch procedure





Figure 4. Block diagram of Iterated Local Search procedure

Hence, in the block diagram in Figure 3, let us comprehend how ILS can be modularized and how its construction can be implemented procedure by procedure alone, as



long we provide a s as input signal in each block. This idea leads us to define the iterated local search algorithm. The algorithm 1 provides a high-level architecture proposed by [STUTZLE and DORIGO 1999].

Algorithm 1 Iterated Local Search

PROCEDURE *Iterated Local Search* $s_0 \leftarrow \text{GenerateInitialSolution}()$ $s^* \leftarrow \text{LocalSearch}(s_0)$ **REPEAT** $s' \leftarrow \text{Perturbation}(s^*, \text{history})$ $s^{*'} \leftarrow \text{LocalSearch}(s')$ $s^* \leftarrow \text{AcceptanceCriterion}(s^*, s^{*'}, \text{history})$ **UNTIL** termination condition met**END**


The procure *GenerateInitialSolution()* runs only once and provides the initial solution s_0 , so that can be the input for the *LocalSearch()* procedure. Then, a *Perturbation()* is applied in s^* and gives a new input with other basin attraction in space S , providing a s' which is the new input input for another call of *LocalSearch()*. The *AcceptanceCriterion()* procedure defines which solution is better. To provide a feedback structure each time, the solution is set to s^* , and the algorithm repeats the last steps until a defined condition is met. These four components, i.e., *GenerateInitialSolution()*, *LocalSearch()*, *Perturbation()* and *AcceptanceCriterion()*, performing together, implements the ILS metaheuristic.

5. Greedy Randomized Adaptive Search

The *Greedy Randomized Adaptive Search Procedures* (GRASP) is also a sophisticated metaheuristics algorithm designed to solve optimization problems. Operating by iteratively constructing and locally searching the solutions and comparing them to provide an acceptable one. The GRASP, in its iterative processes, has the mentioned two main multi-start phases, the *construction* and *local search*. The first phase builds a feasible solution, and the second investigates the neighborhood. Again, the best solution (the *best* is determined by desired solution conditions) is chosen as the result.

The entrusted procedure of the construction phase is implemented as the *GreedyRandomizedConstruction*. Given an initial solution s_0 , the procedure evaluates a set of candidate elements formed by all with the potential to be incorporated into a partial solution s' . The evaluation is obtained using a *greedy* evaluation procedure. The incremental cost usually biases this greedy procedure due to incorporating this candidate element into the partial solution.

Hence, a structure with a restricted candidate list (RCL) formed by the best elements, i.e., those whose incorporation into the partial solution produces the minimum cost, is necessary; this guarantees the greedy facet of the algorithm. The feasibility of an element being definitively incorporated in the partial solution is decided based on a procedure that randomly selects one from the RCL; this guarantees the probabilistic facet of the algorithm. Since the element is set as the next element of the partial solution, the



update of the RCL occurs, and the incremental costs are reevaluated; this guarantees the adaptive facet of the algorithm. Figure ?? illustrates the behavior of this phase.

Those ideas lead us to define the Greedy Randomized Construction Procedure algorithm. The algorithm 2 provides a high-level architecture proposed by [].

Algorithm 2 Greedy Randomized Construction

```
PROCEDURE GreedyRandomizedConstruction( $S_0$ )  
   $s \leftarrow \emptyset$   
  Evaluate the incremental costs of the candidate elements  
  REPEAT  
    Build the restricted candidate list (RCL)  
    Select an element  $s_r$  from the RCL at random  
     $s \leftarrow s \cup s_r$   
    Reevaluate the incremental costs  
  UNTIL Solution is not complete  
  RETURN  $s$   
END
```

The procedure initializes with an empty solution s ; then, each element is evaluated to compose the RCL. Iteratively, the RCL is constructed; then, only one element is chosen, at random, from the RCL and incorporated into the partial solution s . Again, each element is evaluated to compose the new RCL. Those last steps iterates until the partial solution s is completed.


At this, the solutions generated by the *GreedyRandomizedConstruction* are not necessarily optimal. As a subsequential step, improving the current solution is a possibility. The current solution is a neighborhood structure, so applying a *LocalSearch* procedure is possible. This second phase should work exactly as the previously implemented procedure in Section 4.

Those ideas lead us to define the Greedy Randomized Adaptive Search Procedure Algorithm. The algorithm 3 provides a high-level architecture proposed by [].

Algorithm 3 Greedy Randomized Adaptive Search Procedures

```
PROCEDURE GRASP  
   $s_0 \leftarrow \text{GenerateInitialSolution}()$   
   $s \leftarrow s_0$   
  REPEAT  
     $s' \leftarrow \text{GreedyRandomizedConstruction}(s)$   
     $s^* \leftarrow \text{LocalSearch}(s')$   
     $s \leftarrow \text{UpdateSolution}(s^*, s)$   
  UNTIL termination condition met  
END
```

The procedure *GenerateInitialSolution()* runs only once and provides the initial solution s_0 , so that can be the input for the first phase *GreedyrandomizedConstruction()*



procedure. That provides a partial solution s' , the new input for *LocalSearch()*. The *UpdateSolution()* procedure defines which solution is better. To provide a feedback structure each time, the solution is set to s , and the algorithm repeats the last steps until a defined condition is met. These four components, i.e., *GenerateInitialSolution()*, *GreedyrandomizedConstruction()*, *LocalSearch()*, and *UpdateSolution()*, performing together, implements the GRASP metaheuristic.

6. Implementing the TSP Using the ILS

This paper implements a sequential and parallel version of the ILS algorithm with $C++$. OpenMP has been used as a multi-core implementation for parallel TSP. The following paragraphs list the pseudocode algorithms.

Algorithm 4 Sequential ILS algorithm for solving the TSP

INPUT: nodes(n_{id} , X, Y), nodesDimension

OUTPUT: bestSolution[], Cost(bestSolution[]), elapsedTime()


```
1: for all nodes( $n_{id}$ , X, Y) do
2:   searchGraph[]  $\leftarrow$  CalculateEclidianDistance(nodes( $n_{id}$ , X, Y), nodesDimension)
3: end for
4: initialTime  $\leftarrow$  GetTime()
5: initialSolution[]  $\leftarrow$  GreedyProcedure(Graph[])
6: bestSolution[]  $\leftarrow$  2OptimizationProcedure(initialSolution[])
7: repeat
8:   perturbedSolution[]  $\leftarrow$  DoubleBridgeMoveProcedure(bestSolution[])
9:   optimizedSolution[]  $\leftarrow$  2OptimizationProcedure(perturbedSolution[])
10:  bestSolution[]  $\leftarrow$  BetterCostProcedure(optimizedSolution[], bestSolution[])
11: until iteration > iterationDimension - 1
12: finalTime  $\leftarrow$  GetTime()
13: return bestSolution[], Cost(bestSolution[]), elapsedTime(initialTime, finalTime)
```

To operate local search algorithms, the algorithm 4 requires a graph data structure for this implementation. Given the graph is complete, an adjacency matrix was chosen as a data structure. To provide a correct input with a list of nodes(n_{id} , X, Y) and a value with the nodes' dimension, a preprocessing step is necessary, where the input is extracted from a text file with this data.

A provided list nodes(n_{id} , X, Y) includes the *id* of a node, which is necessary to identify the vertex of our graph. A cartesian coordinate (X,Y) provides a position, and with the Euclidian method, our procedure calculates the distance between two nodes. A value of the n_{max} where it is the maximum value for the *id* n .

After a correct provided input, the procedure continues:

- Recording the initial time to provide a final elapsed time;
- Generating a greedy tour as an initial solution to allow a start state for our search;
- Optimizing our initial tour with the 2-Opt local search algorithm and set as our best solution;
- Starting the main loop and repeating for iteration value, predefined;



- ▶ Perturbating our best tour, for provides a way to get out of the basin of attraction of the previews best solution;
- ▶ Optimizing our perturbed tour, again with the 2-Opt local search algorithm and set as our intermediate tour, an optimized solution;
- ▶ Getting the less cost tour, between the previous and the best tour, and setting as our current best solution;
- ▶ Iterating the last steps until the predefined iteration value;
- ▶ Recording the final time to provide a final elapsed time;
- ▶ Returns the output records;

After running the metaheuristics algorithm implementation, we should expect as output the best Hamiltonian cycle (tour) for our problem, the cost value of this cycle, and the elapsed time.

Algorithm 5 Parallel ILS algorithm for solving the TSP

INPUT: nodes(n_{id} , X, Y), nodesDimension


OUTPUT: bestSolution[], Cost(bestSolution[]), elapsedTime()

```
1: for all nodes( $n_{id}$ , X, Y) do
2:   searchGraph[]  $\leftarrow$  CalculateEclidianDistance(nodes( $n_{id}$ , X, Y), nodesDimension)
3: end for
4: initialTime  $\leftarrow$  GetTime()
5: initialSolution[]  $\leftarrow$  GreedyProcedure(Graph[])
6: bestSolution[]  $\leftarrow$  2OptimizationProcedure(initialSolution[])
7: parallel // Each thread executes each command.
8: threadSolution[]  $\leftarrow$  bestSolution[]
9: repeat
10:  perturbedSolution[]  $\leftarrow$  DoubleBridgeMoveProcedure(bestSolution[])
11:  optimizedSolution[]  $\leftarrow$  2OptimizationProcedure(perturbedSolution[])
12:  threadSolution[]  $\leftarrow$  BetterCostProcedure(optimizedSolution[], threadSolution[])
13:  threadSolutionWeight  $\leftarrow$  Cost(bestSolution[])
14: until iteration > interationDimension - 1 // Each thread executes  $\frac{interationDimension - 1}{threadsDimension}$ 
15: for each thread
16:  if threadSolutionWeight < bestSolutionWeight then
17:    bestSolution[]  $\leftarrow$  threadSolution[]
18:  end if
19: end parallel
20: finalTime  $\leftarrow$  GetTime()
21: return bestSolution[], Cost(bestSolution[]), elapsedTime(initialTime, finalTime)
```

Given the same problem, the input and output are equal to the sequential ILS, only differing in the algorithm itself.

After a correct provided input, the procedure continues:

- ▶ Recording the initial time to provide a final elapsed time;
- ▶ Generating a greedy tour as an initial solution to allow a start state for our search;
- ▶ Optimizing our initial tour with the 2-Opt local search algorithm and set as our best solution;



- ▶ Starting the parallel region, where each thread executes each step simultaneously unless there is a specific region where each thread executes thread by thread alone;
- ▶ Setting for each thread tour the best initial tour. Each thread operates the procedures only in its particular tour list;
- ▶ Starting the main loop and repeating for each thread $\frac{iterationsDimension}{threadsDimension}$ times, where iterationsDimension is predefined;
- ▶ Perturbating our best tour, for provides a way to get out of the basin of attraction of the previews best solution;
- ▶ Optimizing our perturbed tour, again with the 2-Opt local search algorithm and set as our intermediate tour, an optimized solution;
- ▶ Getting the less cost tour, between the previous and the best tour, and setting as our current best solution;
- ▶ Iterating the last steps until the predefined value for each thread;
- ▶ Setting the best tour between all the threads as the best solution;
- ▶ Recording the final time to provide a final elapsed time;
- ▶ Returns the output records;

7. Implementing the TSP Using the GRASP


This paper implements a sequential and parallel version of the GRASP algorithm with C + +. OpenMP has been used as a multi-core implementation for parallel TSP. The following paragraphs list the pseudocode algorithms.

Algorithm 6 Sequential GRASP algorithm for solving the TSP

INPUT: nodes(n_{id} , X, Y), nodesDimension

OUTPUT: bestSolution[], Cost(bestSolution[]), elapsedTime()

```
1: for all nodes( $n_{id}$ , X, Y) do
2:   searchGraph[]  $\leftarrow$  CalculateEclidianDistance(nodes( $n_{id}$ , X, Y), nodesDimension)
3: end for
4: initialTime  $\leftarrow$  GetTime()
5: initialSolution[]  $\leftarrow$  GreedyProcedure(Graph[])
6: bestSolution[]  $\leftarrow$  initialSolution[]
7: repeat
8:   partialSolution[]  $\leftarrow$  GreedyRandomizedConstruction(bestSolution[])
9:   optimizedSolution[]  $\leftarrow$  2OptimizationProcedure(partialSolution[])
10:  bestSolution[]  $\leftarrow$  BetterCostProcedure(optimizedSolution[], bestSolution[])
11: until iteration > iterationDimension - 1
12: finalTime  $\leftarrow$  GetTime()
13: return bestSolution[], Cost(bestSolution[]), elapsedTime(initialTime, finalTime)
```



Algorithm 7 Greedy randomized construction algorithm for solving the TSP

INPUT: solution[], nodesDimension, alpha

OUTPUT: newSolution[]

```
1: currentSolutionNode ← randomNode(solution[])
2: initialTime ← GetTime()
3: newSolution[] ← currentSolutionNode
4: repeat
5:   minCost ← getMinCost(solution[])
6:   maxCost ← getMaxCost(solution[])
7:   feasibleCost ← minCost + (alpha * (maxCost - minCost))
8:   for  $s$  from 1 to nodesDimension do
9:     if (newSolution[]  $\cup$  solution[ $s$ ]  $\leq$  feasibleCost) then
10:      recentCandidateList[] ← solution[ $s$ ]
11:    end if
12:  end for
13:  currentSolutionNode ← randomNode(recentCandidateList[])
14:  newSolution[] ← currentSolutionNode
15: until iteration > nodesDimension - 1
16: return newSolution[]
```


To operate local search algorithms, the algorithm 8 requires a graph data structure for this implementation. Given the graph is complete, an adjacency matrix was chosen as a data structure. To provide a correct input with a list of nodes(n_{id} , X, Y) and a value with the nodes' dimension, a preprocessing step is necessary, where the input is extracted from a text file with this data.

A provided list nodes(n_{id} , X, Y) includes the id of a node, which is necessary to identify the vertex of our graph. A cartesian coordinate (X,Y) provides a position, and with the Euclidian method, our procedure calculates the distance between two nodes. A value of the n_{max} where it is the maximum value for the id n .

After a correct provided input, the procedure continues:

- ▶ Recording the initial time to provide a final elapsed time;
- ▶ Generating a greedy tour as an initial solution to allow a start state for the *GreedyRandomizedConstruction*;
- ▶ Setting our initial solution as our best solution at this step;
- ▶ Starting the main loop and repeating for iteration value, predefined;
- ▶ Constructing our greedy randomized partial, for the phase one;
- ▶ Optimizing our partial tour with the 2-Opt local search algorithm and set as our optimized solution, for the phase two;
- ▶ Getting the less cost tour, between the previous and the best tour, and setting as our current best solution;
- ▶ Iterating the last steps until the predefined iteration value;
- ▶ Recording the final time to provide a final elapsed time;
- ▶ Returns the output records;

After running the metaheuristics algorithm implementation, we should expect as output the best Hamiltonian cycle (tour) for our problem, the cost value of this cycle, and the elapsed time.



Algorithm 8 Sequential GRASP algorithm for solving the TSP

INPUT: nodes(n_{id} , X, Y), nodesDimension

OUTPUT: bestSolution[], Cost(bestSolution[]), elapsedTime()

```
1: for all nodes( $n_{id}$ , X, Y) do
2:   searchGraph[]  $\leftarrow$  CalculateEclidianDistance(nodes( $n_{id}$ , X, Y), nodesDimension)
3: end for
4: initialTime  $\leftarrow$  GetTime()
5: parallel           // Each thread executes each command.
6: initialSolution[]  $\leftarrow$  GreedyProcedure(Graph[])
7: bestSolution[]  $\leftarrow$  initialSolution[]
8: repeat
9:   partialSolution[]  $\leftarrow$  GreedyRandomizedConstruction(bestSolution[])
10:  optimizedSolution[]  $\leftarrow$  2OptimizationProcedure(partialSolution[])
11:  bestSolution[]  $\leftarrow$  BetterCostProcedure(optimizedSolution[], bestSolution[])
12: until iteration > iterationDimension - 1 // Each thread executes  $\frac{\text{iterationDimension} - 1}{\text{threadsDimension}}$ .
13: for each thread
14:   if threadSolutionWeight < bestSolutionWeight then
15:     bestSolution[]  $\leftarrow$  threadSolution[]
16:   end if
17: end parallel
18: finalTime  $\leftarrow$  GetTime()
19: return bestSolution[], Cost(bestSolution[]), elapsedTime(initialTime, finalTime)
```

Given the same problem, the input and output are equal to the sequential GRASP, only differing in the algorithm.

After a correct provided input, the procedure continues:

- ▶ Recording the initial time to provide a final elapsed time;
- ▶ Starting the parallel region, where each thread executes each step simultaneously unless there is a specific region where each thread executes thread by thread alone;
- ▶ Generating a greedy tour as an initial solution to allow a start state for the *GreedyRandomizedConstruction*;
- ▶ Setting our initial solution as our best solution at this step;
- ▶ Starting the main loop and repeating for iteration value, predefined;
- ▶ Constructing our greedy randomized partial for the phase one;
- ▶ Optimizing our partial tour with the 2-Opt local search algorithm and set as our optimized solution, for the phase two;
- ▶ Getting the less cost tour, between the previous and the best tour, and setting as our current best solution;
- ▶ Iterating the last steps until the predefined value for each thread;
- ▶ Setting the best tour between all the threads as the best solution;
- ▶ Recording the final time to provide a final elapsed time;
- ▶ Returns the output records;

8. Experiment Setup

8.1. Implement Language and Frameworks

The implementation uses C++20 and OpenMP 5.2, compiled using GCC 12.2.0.

8.2. Testing Enviroment

The code was compiled and run on a machine x86_64, AMD Ryzen 5 5600GT with Radeon Graphics, 6 cores, and 12 threads. With Debian 12.2.0-14. The program was compiled using `g++ component.cpp node.cpp scanner.cpp functions.cpp mainExec.cpp -o ../bin/TspPar -fopenmp -Wall -pedantic`, and tested using a shell script.

8.3. Datasets - Inputs

For the datasets, nine TSP benchmarks were downloaded from TSPLIB [TspLib], which includes d198, a280, lin318, pcb442, rat783, u1060, pcb1173, d1291, and fl1577 have been used in this write-up for evaluating the proposed algorithm performance. The selected benchmarks varied in several cities where each city was represented by a 2D-Euclidian coordinate.

The table 1 displays benchmarks properties.

Benchmark name	Dimension	Description
d198.tsp	198 city	Drilling problem (Reinelt)
a280.tsp	280 city	drilling problem (Ludwig)
lin318.tsp	318 city	318-city problem (Lin/Kernighan)
pcb442.tsp	442 city	Drilling problem (Groetschel/Juenger/Reinelt)
rat783.tsp	783 city	Rattled grid (Pulleyblank)
u1060.tsp	1060 city	Drilling problem problem (Reinelt)
pcb1173.tsp	1173 city	Drilling problem (Juenger/Reinelt)
d1291.tsp	1291 city	Drilling problem (Reinelt)
fl1577.tsp	1577 city	Drilling problem (Reinelt)

Table 1. TSP benchmarks' properties

8.4. Evaluation Measures

This section presents the evaluation metrics used to evaluate the proposed method.

1. Distance: the value of the best route found.
2. Accuracy: the percentage of retrieving the best route correctly. The optimal solution is provided in TSPLIB [TspLib].
3. Running time: is the duration between the end and the beginning of the main part (ILS algorithm itself) of the program running, using `chrono::high_resolution_clock::now()` for the sequential program and `omp_get_wtime()` for the parallel program.

$$RT = \text{finish_time} - \text{start_time} \quad (3)$$

4. Speedup: is the improvement in speed of execution of a task executed on two similar architectures with different resources:

$$SP = \frac{T_s}{T_p(n)} \quad (4)$$

► T_s : Sequential running time;

- ▶ $T_p(n)$: Parallel running time in function of n ;
- ▶ n : Number of cores.

5. Efficiency: represents the speedup divided by the number of cores

$$EF = \frac{SP}{n} \quad (5)$$

9. Experimental results

The results are computed over an average of 30 trials for each instance, and those results are presented in Tables ??

Table 2. Average distance of each symmetric TSP benchmark, the Sequential and Parallel ILS, then the optimal solution.

Instance	Sequential Distance	Parallel Distance	Optimal
d198	1.628×10^4	1.721×10^4	1.578×10^4
a280	2.643×10^3	2.729×10^3	2.579×10^3
lin318	4.382×10^4	4.763×10^4	4.203×10^4
pcb442	5.172×10^4	5.348×10^4	5.078×10^4
rat783	9.271×10^3	9.659×10^3	8.806×10^3
u1060	2.467×10^5	2.416×10^5	2.241×10^5
pcb1173	6.066×10^4	6.498×10^4	5.689×10^4
d1291	—	5.690×10^4	5.080×10^4
fl1577	2.342×10^4	2.966×10^4	2.2249×10^4

Table 3. Accuracy of each symmetric TSP benchmark for Sequential and Parallel ILS.

Instance	Sequential Accuracy (%)	Parallel Accuracy (%)
d198	96.91	91.97
a280	97.58	94.50
lin318	95.91	88.25
pcb442	98.18	94.95
rat783	94.98	91.17
u1060	90.84	92.75
pcb1173	93.79	87.55
fl1577	95.00	75.02

Table 4. Running time, speedup, and efficiency of each symmetric TSP benchmark for Sequential and Parallel ILS.

Instance	Sequential Time (s)	Parallel Time (s)	Speedup	Efficiency (%)
d198	6.114×10^1	8.649	7.069	58.91
a280	1.755×10^2	2.411×10^1	7.279	60.66
lin318	2.775×10^2	3.941×10^1	7.041	58.68
pcb442	7.463×10^2	1.018×10^2	7.331	61.09
rat783	4.620×10^3	5.999×10^2	7.701	64.18
pcb1173	1.676×10^4	2.183×10^3	7.679	63.99
fl1577	4.293×10^4	5.797×10^3	7.406	61.71

10. Conclusão

References

- Bondy, J. A. and Murty, U. S. (1976). *Graphy Theory With Applications*. Elsevier Science Publishing Co., Inc., 1^o edition.
- Esra'a Alhenawi, Ruba Abu Khurma, R. D. and Hussien, A. G. (2024). Solving traveling salesman problem using parallel river formation dynamics optimization algorithm on multi-core architecture using apache spark. *International Journal of Computational Intelligence Systems*.
- Helena R. Lorengo, O. M. and STUTZLE, T. Iterated local search.
- Pan-Li Zhang, Xiao-Bo Sun, J.-Q. W. H.-H. S. J.-L. B. and Zhang, H.-Y. (2022). The discrete carnivorous plant algorithm with similarity elimination applied to the traveling salesman problem. *mathematics*.
- Reinelt, G. (1994). *The Traveling Salesman: Computational Solutions for TSP Applications*. Springer-Verlag Berlin Heidelberg, 1^o edition.
- Scianna, M. (2024). The addaco: A bio-inspired modified version of the ant colony optimization algorithm to solve travel salesman problems. *Mathematics and Computers in Simulation*.
- STUTZLE, T. and DORIGO, M. (1999). Aco algorithms for the traveling salesman problem.
- Thomas H. Cormen, Charles E. Leiserson, R. L. R. and Stein, C. (2009). *Introduction to Algorithms*. The MIT Press, 3^o edition.
- TspLib. Disponível em: <http://comopt.ifl.uni-heidelberg.de/software/TSPLIB95/>. Acesso em: 11 de setembrbo 2023.
- Xing Li, S. Z. and Shao, P. (2024). Discrete artificial bee colony algorithm with fixed neighborhood search for traveling salesman problem. *Engineering Applications of Artificial Intelligence*.